

001

002

003

The Text Beside the Image: Detection, Utility and
Leakage for Trustworthy Multimodal Medical
Data and Beyond

001

002

003

004

Anonymous ACCV 2026 Submission

004

005

Paper ID #*****

005

006

007

008

009

010

011

012

013

014

015

016

017

018

019

020

021

022

023

Abstract. Medical images are released with the free text that describes them, and protecting the image does not protect the report. We measure identifier detection, downstream utility and residual identity leakage on the same documents, with the pseudonymisation policy as the variable under test: 15 detectors, three release conditions and four corpora of clinical letters, legal judgments, news and other genres, and e-mail, in German, English, Chinese and Arabic. A fixed 13-detector union replaces almost every person mention in the clinical letters and 0.890 of them in e-mail. Frequency matching with a public name list then recovers no identity at all; every name it returns is one the detector missed and the release printed in clear text. Cross-document linkage reaches 0.0095 without training and 0.0407 with it on e-mail, against 1/3816 chance and a 0.706 ceiling on unmodified text. Replacement form bounds neither attack: the typed-placeholder release is a character-for-character rewrite of the surrogate release, so surrogates cannot protect against an adversary who knows the scheme. What remains is governed by identifiers the detector missed, and has to be reported per document, per case and per language.

006

007

008

009

010

011

012

013

014

015

016

017

018

019

020

021

022

023

024

025

Keywords: pseudonymisation · clinical text · multimodal medical data · privacy evaluation

024

025

026

1 Introduction

026

027

028

029

030

031

032

033

034

035

036

Clinical data are multimodal: a chest radiograph travels with its report, a cardiology study with its discharge letter. Releasing either half requires both to be protected, since an image may need defacing or removal of burned-in text while its report needs identifier replacement. Artificial intelligence scales both tasks, and its errors vary by distribution, language, script and identifier class. Packhäuser et al. [18] re-identified patients from de-identified chest radiographs and priced a defence against diagnostic utility [19]; pathological speech has been evaluated against privacy, utility and fairness [23]; clinical text is instead kept with the model inside the hospital [9]. We ask what survives transformation of the text, and what protecting it costs.

027

028

029

030

031

032

033

034

035

036

De-identification reduces identifying information. Reversible *pseudonymisation* replaces an identifier with a realistic *surrogate*; *anonymisation* prevents singling out. Stable pseudonyms preserve longitudinal utility but hand an adversary a persistent handle. European guidance treats keyed message authentication codes as robust [8], and the health-informatics standard is revised accordingly [6].

Detection benchmarks ask whether an identifying span was found [10, 25, 27]; utility studies measure downstream change [1, 26] or information loss [24]; re-identification studies measure attack success. Frameworks combine technical performance, information loss and an intruder test [12], while surrogate systems and masking benchmarks cover single pieces [7, 17, 21]. None states together what one policy buys and what it costs. Data-protection officers release documents and cases, not average tokens, and one unchanged mention can identify someone after thousands were removed, so we measure surviving findings, exposed cases and linkable entities rather than token rates alone.

One design covers four corpora in German, English, Chinese and Arabic: clinical letters, legal judgments, news and other genres, and e-mail. Every document receives the same detector comparison, release conditions and measurements. Two predictions are put at risk: that detector precision governs utility more strongly than the choice between surrogates and typed placeholders, and that detection sensitivity moves the two attacks in opposite directions, because removing a name obstructs context linkage while leaving a larger, more stable population for frequency matching. The clinical corpus has a known weakness: its gold names were inserted into an already de-identified release, so its name distribution is synthetic.

2 Material and Methods

2.1 Corpora

Reference provenance differs, so detection denominators are not interchangeable (Table 1). TAB manually annotates identifiers and document-scoped co-reference in court judgments [21]; OntoNotes covers three languages but was annotated for linguistics rather than privacy. A *named entity* is a predefined person, organisation, location or other type, its occurrences *mentions* forming a co-reference chain. Enron derives genuine cross-document identity from message headers but contains exposed personal information [16]. CARDIO:DE is the clinical anchor, its gold identifiers inserted into an already de-identified release, which makes its name distribution synthetic.

One denominator follows. CARDIO:DE gold covers only inserted identifiers, and 84.7% of its gold tokens are dates, which the detectors do find but the conditions pass through by design. Sensitivity, replacement, exposure and the attacks are therefore all scored over the types actually replaced, and the retained types are in none of those denominators. Retained dates are themselves a re-identification channel, and a release keeping them needs date shifting. A TAB applicant is usually named once but has a median of 17 later role mentions, and only Enron has natural cross-document identity.

Table 1: The corpora separate domain, language, and where the reference annotation came from.

Corpus	Docs	Tokens	Ment./doc	Lang.	Reference annotation
CARDIO:DE	400	885 059	78.3	de	constructed
TAB	1268	1 844 804	81.2	en	manual
OntoNotes	5994	4 726 594	45.5	en/zh/ar	linguistic
Enron	58 636	15 507 925	14.4	en	header-derived




Fig. 1: One set of detector spans generates three release conditions, on an invented header line in the style of a discharge letter. Detected identifiers (teal) become surrogates in B and typed placeholders in C; dates (red) pass through by design. B and C replace the same spans, so C is a rewrite of B. Every condition meets every measurement family on the same documents.

2.2 Conditions

Condition A is unmodified text. In B, each detected identifier becomes a realistic *surrogate*, such as another name, stable across the corpus, indexed by HMAC-SHA256, a keyed one-way digest irreproducible without the secret. Condition C uses an unindexed placeholder such as [PERSON]. B and C differ only in replacement form, following European guidance and prior e-mail pseudonymisation [7, 8], and replace identical spans.

Before hashing, the key normaliser casefolds mentions, collapses whitespace and removes titles and punctuation, so *Dr. Weber* matches *weber* but *Weber* does not match *Kay Weber*. Unrelated people may therefore share a key while one person’s initials and full name may not; Sect. 3.3 reports both failures. Figure 1 evaluates every condition on identical documents and measurements.

2.3 Detectors, Ensembles and Operating Points

The 15 frozen detectors comprise rules, zero-shot and fine-tuned encoders, multilingual and German clinical recognisers, and eight gateway language models; the repository lists every model, version and prompt. Cached spans are combined by union, intersection or at-least- k vote. Two rules recur below: the *permissive* rule is the union of all 15, the *precise* rule a majority, 8 of the 15. Leakage uses the configuration we recommend, a fixed 13-detector union; two large language models remain available singly but are excluded from it, because preliminary runs made them one to two orders of magnitude slower per document without

improving detection. A proportional output limit truncated 92% of Chinese replies, which we corrected and re-ran rather than score truncation as a detection failure.

Each corpus receives fast and maximum points for sensitivity, the share of replaced-type identifier tokens detected, and specificity, the share of all other tokens left untouched; an ensemble run in parallel costs its slowest member. Neither rate is readable alone, so both are given at every point. A maximum point is non-dominated, the fast point the cheapest within 10% of it, subject to a sensitivity floor that stops an empty detector winning on specificity. This search covers singles, pairs and triples, so the recommended 13-detector ensemble lies outside it and may exceed the maximum of Table 2.

2.4 Measurements

Detection is scored by token and by entity, an entity counting as protected only when every one of its mentions was replaced. Information-weighted precision scales each masked token by how poorly its context predicts it, so removing a predictable word costs less than removing an informative one, and exposure uses document, case and entity denominators, because few leaked tokens can expose many people. Frozen models analyse all three conditions for named-entity agreement, medication extraction and section classification; agreement is span F1 between the outputs on original and transformed text, and the original-text score accompanies every comparison. Continuous pairs use the Wilcoxon signed-rank test, binary correctness McNemar’s test, and Benjamini–Hochberg correction applies within task families. Utility and attack figures carry one standard deviation over documents, detector combinations or folds; detection and exposure are exact population proportions.

A *collision* gives distinct people one surrogate, *fragmentation* gives one person several within a document, and *drift* changes it between documents. These apply to B, which has a mapping, and not to placeholders, which merge every entity of one type by design.

2.5 The Attacks

Frequency matching aligns surrogate frequencies with real-name frequencies. Published attacks on pseudonymised free text target the identifiers a system missed rather than the surrogate distribution it produced, as the parrot attack does [4], so we transfer the alignment model from attacks on encoded identifiers in record linkage and disclosure-risk work [14, 29]. Two priors run over one observation, following Kuzu et al., Niedermeyer et al. and Vatsalan et al. [11, 15, 28]: a corpus-internal prior granting perfect knowledge, an oracle upper bound, and a public name-frequency list, the realistic adversary.

The attacker tags the released text itself and so also sees names the defender missed. Every detected surface stays in the denominator, including ambiguous and unresolved bins [5], and if k entities share a surrogate a recovery scores $1/k$ rather than being deleted. Abstention occurs where $(\max - \max_2)/\sigma < \varphi$,

Table 2: Operating points on both error rates. Sensitivity is over the identifier types the conditions replace, specificity over every other token. Cost is seconds per document; the rule is $\cup$ union, $\cap$ intersection, v at-least- k vote, s a single detector.

Operating point	CARDIO:DE			TAB			OntoNotes			Enron		
	Sens.	Spec.	Cost	Sens.	Spec.	Cost	Sens.	Spec.	Cost	Sens.	Spec.	Cost
max. sensitivity	0.992	0.9670	57.633 ^{$\cup$}	0.925	0.8986	14.210 ^{$\cup$}	0.747	0.9763	9.450 ^{$\cup$}	0.973	0.7720	25.557 ^{$\cup$}
fast sensitivity	0.895	0.9769	0.301 ^{$\cup$}	0.893	0.9242	0.219 ^{$\cup$}	0.684	0.9731	5.376 ^{$\cup$}	0.899	0.7419	0.057 ^{$\cup$}
max. specificity	0.523	0.9997	116.875 ^{v}	0.501	0.9959	98.110 ^{v}	0.503	0.9960	0.103 ^{$\cup$}	0.588	0.9428	20.931 ^{$\cap$}
fast specificity	0.781	0.9862	0.263 ^{s}	0.520	0.9934	0.117 ^{s}	0.503	0.9960	0.103 ^{$\cup$}	0.674	0.8787	0.026 ^{s}

computed from the attacker’s own scores [13]. We report the answered-query rate and $c@1$, which credits warranted abstention [20], beside a chance rate from the candidate set [22]. Realistic surrogates exist so residual identifiers are hard to pick out [3], which is what the public prior tests; sweeping its size and vintage measures what better knowledge is worth [2].

Context linkage ranks identities by cosine similarity between surrounding-word profiles; *learned linkage* fits the same task on labelled, entity-disjoint examples. Both use identical document-disjoint folds, and only the labels of held-out queries are hidden: those people stay in the complete candidate list with every distractor, so neither attack is told that a query has a match, and their difference measures learning rather than gallery size. Cross-document identity permits this on Enron and CARDIO:DE but not on TAB or OntoNotes.

LLM candidate ranking gives a language model one marked mention and ten candidate names, with or without auxiliary text about each candidate from the released condition. Because B and C replace identical spans, B can be rewritten into C, so we report this attack on B as presented, on B after that rewrite, and against a ranker told the scheme. The one-in-ten closed world measures an attacker holding a shortlist, not population re-identification. A *dictionary lookup* would hash candidate names; it is defined but not run, because no condition here is unkeyed.

3 Results

3.1 Detection

Table 2 gives the cost–quality front. The winning configuration differs by corpus, so it is a set of four answers rather than one transferable recipe.

The fast point is usually not the worse one: on CARDIO:DE it gives up 0.097 sensitivity but is more specific than the maximum, 0.9769 against 0.9670, at 191 times less cost, and Enron repeats this at 0.057 s against 25.557 s per message. The two OntoNotes specificity rows are one ensemble, because the 10 % band admits all 414 candidates.

Language is a separate axis, and how separate depends on what is detected. Under the recommended ensemble, sensitivity is 0.989 on English, 0.868 on Chinese and 0.614 on Arabic, and the pooled 0.801 describes none of them: Arabic

Table 3: Leakage from the recommended 13-detector union; all columns concern people. *Sens.* is person sensitivity and *Spec.* the share of non-identifier tokens left alone: neither is readable without the other. *Exp.* is the share of people keeping a mention in clear text, *Frequency* the identities named from frequency alone with a public list and, bracketed, with the corpus’s own distribution (an upper bound, not a risk). *Linkage* is Rank-1 over the complete candidate list against chance 1/gallery.

Corpus	Sens.	Spec.	Exp. (%)	Frequency	Context link.	Learned link.	Chance
CARDIO:DE	1.000	0.869	0.1	0 (1)	0.0000 ± 0.0000	0.0071 ± 0.0160	1/207
Enron	0.890	0.624	18.0	0 (0)	0.0095 ± 0.0012	0.0407 ± 0.0029	1/3816
TAB	0.996	0.850	0.8	0 (1)	not measurable		–
OntoNotes	0.935	0.932	4.9	0 (1)	not measurable		–

supplies 45 % of the identifier tokens from 7 % of the documents, dominating the pooled score while being the hardest case. People transfer far better, at 0.991, 0.932 and 0.878, pooling to 0.935, so the thirty-seven-point spread belongs to the other identifier classes rather than to finding people, which is what the attacks try to do.

3.2 Utility

Detector precision governs utility more strongly than the replacement form. Named-entity agreement on CARDIO:DE is 0.611 ± 0.082 with surrogates under the permissive rule and 0.973 ± 0.019 under the precise rule; with placeholders it is 0.004 ± 0.008 and 0.674 ± 0.082 . The placeholder figures measure the metric as much as the release, because the frozen recogniser never tags [PERSON] as a name, so every replaced mention counts as a disagreement and the comparison is informative between rules rather than between forms.

The other two tasks carry their original-text scores and admit the paired test. Medication extraction scores 0.371 ± 0.182 on unmodified text; under the precise rule it moves to 0.364 ± 0.189 with surrogates, a paired difference of -0.008 ± 0.178 not distinguishable from zero ($q = 0.80$), and under the permissive rule it falls to 0.346 ± 0.181 , a difference of -0.020 ± 0.182 that is ($q = 5.6 \times 10^{-3}$). Section classification scores 0.749 ± 0.093 on unmodified text and loses 0.072 ± 0.086 to surrogates against 0.038 ± 0.089 to placeholders, both significant ($q = 1.6 \times 10^{-37}$, $q = 2.6 \times 10^{-14}$; rank-biserial -0.846 , -0.531). At document level the realistic surrogate is therefore not merely unnecessary but worse than the typed placeholder, while at token level it is the only form that survives.

3.3 Leakage

Table 3 reports the recommended 13-detector union, the intended release configuration.

Frequency matching recovers nothing. With a public name list it names no identity on any corpus; every name it returns is one the detector missed

and the release printed in clear text, which measures the defender’s sensitivity rather than the adversary’s inference. Even the corpus’s own distribution, which no real attacker holds, aligns at most one identity, and raising the abstention threshold from $\varphi = 0$ to $\varphi = 0.05$ makes the attack decline all but 11 of 25,677 CARDIO:DE queries and all of them elsewhere, so the alignment’s $c@1$ is zero at every threshold. This holds for one attack and these priors and does not establish safety against others.

Linkage succeeds where identity is real, and needs two anchors to be read. On Enron the fixed similarity recovers 0.706 of held-out people from unmodified text; after the recommended release it recovers 0.0095 ± 0.0012 , still 36 times the $1/3816$ chance rate but two orders of magnitude below that ceiling. Training the metric on entity-disjoint examples raises it to 0.0407 ± 0.0029 , four times the fixed attack and 155 times chance, so the residual risk belongs to an adversary who can learn. On CARDIO:DE, with 207 identities in the gallery, the fixed attack recovers nothing and the trained one 0.0071 ± 0.0160 against chance 0.0048, not distinguishable from guessing. Recall bounds this: two detectors agreeing lowers Enron person sensitivity to 0.732, raises fixed linkage fivefold to 0.0500 and takes the names in clear text from 180 to 1,868, while intersection leaves person sensitivity at 0.037 and linkage at 0.2380.

The third attack shows why the replacement form cannot be credited with a difference. Ranking one marked mention against ten candidates, a language model falls from 0.995 on unmodified TAB text to 0.110 with surrogates, near the 0.100 chance rate, but reaches 0.340 with typed placeholders; CARDIO:DE (0.085 against 0.120) and OntoNotes (0.290 against 0.630) order the same way. Read as protection this says surrogates defend better, and that reading is wrong. The conditions replace identical spans, so overwriting each replaced span of the surrogate release with [PERSON] reproduces the placeholder release character for character, verified on all 1,268 TAB and all 5,994 OntoNotes documents. The placeholder release is therefore computable from the surrogate release by anyone able to locate those spans, and this is measured rather than argued: an adversary merely told the scheme, with the text left untouched, reaches 0.165, and one neutralising name-like spans with a public name list reaches 0.145, against the 0.340 of the rewrite itself. The surrogate buys the credulity of a ranker that takes it at face value, not protection, and both forms carry the higher risk.

Sensitivity hides exposure, and the denominator decides what a number means. The recommended release leaves one CARDIO:DE person mention of 4,396 in clear text, one entity of 1,957 and one patient of 270; over every identifier type, 8.5 % of letters carry something and 12.2 % of patients do. TAB leaves a name in 44 of 1,268 judgments and OntoNotes in 736 of 5,994 documents, while Enron leaves one in 18.0 % of its people and something in 45.8 % of its messages.

The key normaliser both merges and splits identities: on Enron, ORGANISATION surrogates collide at 17.3 % and drift at 48.6 %, and PERSON drift reaches 37.7 % against 5.0 % on CARDIO:DE, whose lower figure reflects constructed recurrence with fewer natural variants rather than generalisation.

4 Discussion

Four recommendations follow, and one prediction does not survive.

Prefer a large union to a consensus rule, and price it. Agreement buys specificity, 0.869 against 0.986 on CARDIO:DE, and pays in sensitivity, which every leakage measure follows: two votes on Enron take person sensitivity from 0.890 to 0.732 and raise fixed linkage fivefold. **State the attacker’s knowledge with the number.** A corpus-internal prior is an upper bound, not a risk, and here the two differ by the whole result. **Report the denominator you release on.** No corpus-level score describes a document, a case or a language.

Choose the replacement form for utility, not for privacy. The placeholder release is a rewrite of the surrogate release, so neither is safer against an adversary who knows the scheme, and merely telling the ranker so lifts it from 0.110 to 0.165 on unchanged text. The surrogate’s real advantage is that token-level tasks survive it and do not survive [PERSON]; at document level the placeholder is better.

The prediction that sensitivity moves the two attacks in opposite directions is refuted. Better detection was expected to starve context linkage while leaving a larger population for frequency alignment. Instead alignment recovered nothing at any sensitivity and both linkage attacks fell as it rose: both are bounded by the identifiers the detector missed.

The limits are specific. Linkage is measurable only on Enron and CARDIO:DE, because TAB and OntoNotes link mentions within a document; people recur there, so those figures understate rather than bound the risk. CARDIO:DE names were inserted into an already de-identified release, making its frequency distribution synthetic and its clinical realism partial. Dates pass through by design, a channel this pipeline does not close. The candidate-ranking attack presents ten names, bounding an adversary with a shortlist rather than one searching a population. Every number describes one release under one key.

5 Conclusion

Safe sharing of clinical data needs both modalities protected: a guarded image does not compensate for a report that names the patient, and a guarded report does not remove identity from the image. Artificial intelligence makes both possible at scale, but its failures must be measured where data are released. For text we priced detection, utility and re-identification on the same documents. A large union ensemble replaces almost every person mention and holds both attacks near their chance rates, and what remains is bounded by the identifiers the detector missed rather than by the replacement scheme, which is one rewrite away from its alternative. Detector precision governs utility more than the choice between surrogates and placeholders, and sensitivity understates exposure per document, case and language. There is therefore no single anonymisation score: a release should state its operating point, per-language errors, surviving cases, attacker knowledge and stability. Our operating points and measurements are released with it.

References

1. Berg, H., Henriksson, A., Dalianis, H.: The impact of de-identification on downstream named entity recognition in clinical text. In: Proceedings of the 11th International Workshop on Health Text Mining and Information Analysis. pp. 1–11. Association for Computational Linguistics (2020). <https://doi.org/10.18653/v1/2020.louhi-1.1> 2
2. Bindschadler, V., Grubbs, P., Cash, D., Ristenpart, T., Shmatikov, V.: The tao of inference in privacy-protected databases. Proceedings of the VLDB Endowment **11**(11), 1715–1728 (2018). <https://doi.org/10.14778/3236187.3236217> 5
3. Carrell, D., Malin, B., Aberdeen, J., Bayer, S., Clark, C., Wellner, B., Hirschman, L.: Hiding in plain sight: use of realistic surrogates to reduce exposure of protected health information in clinical text. Journal of the American Medical Informatics Association **20**(2), 342–348 (2013). <https://doi.org/10.1136/amiajnl-2012-001034> 5
4. Carrell, D.S., Cronkite, D.J., Li, M.R., Nyemba, S., Malin, B.A., Aberdeen, J.S., Hirschman, L.: The machine giveth and the machine taketh away: a parrot attack on clinical text deidentified with hiding in plain sight. Journal of the American Medical Informatics Association **26**(12), 1536–1544 (2019). <https://doi.org/10.1093/jamia/ocz114> 4
5. Dankar, F.K., El Emam, K., Neisa, A., Roffey, T.: Estimating the re-identification risk of clinical data sets. BMC Medical Informatics and Decision Making **12**(1) (2012). <https://doi.org/10.1186/1472-6947-12-66> 4
6. DIN Deutsches Institut für Normung e. V.: DIN EN ISO 25237:2026-06, medizinische informatik — pseudonymisierung (ISO/DIS 25237:2026); deutsche und englische fassung prEN ISO 25237:2026. Tech. rep., DIN Media GmbH (2026). <https://doi.org/10.31030/3696721> 2
7. Eder, E., Krieg-Holz, U., Hahn, U.: De-identification of emails: Pseudonymizing privacy-sensitive data in a German email corpus. In: Proceedings of the International Conference on Recent Advances in Natural Language Processing (RANLP 2019). pp. 259–269. INCOMA Ltd., Varna, Bulgaria (Sep 2019). https://doi.org/10.26615/978-954-452-056-4_030 2, 3
8. European Union Agency for Cybersecurity: Data pseudonymisation: Advanced techniques and use cases. Tech. rep., ENISA (2021) 2, 3
9. Hou, Y., Bert, C., Gomaa, A., Lahmer, G., Höfler, D., Weissmann, T., Voigt, R., Schubert, P., Schmitter, C., Depardon, A., Semrau, S., Maier, A., Fietkau, R., Huang, Y., Putz, F.: Fine-tuning a local llama-3 large language model for automated privacy-preserving physician letter generation in radiation oncology. Frontiers in Artificial Intelligence **7**, 1493716 (2025). <https://doi.org/10.3389/frai.2024.1493716> 1
10. Jha, P.: PIIBench: A unified multi-source benchmark corpus for personally identifiable information detection (2026). <https://doi.org/10.5281/zenodo.19533263> 2
11. Kuzu, M., Kantarcioglu, M., Durham, E.A., Toth, C., Malin, B.: A practical approach to achieve private medical record linkage in light of public resources. Journal of the American Medical Informatics Association **20**(2), 285–292 (2013). <https://doi.org/10.1136/amiajnl-2012-000917> 4
12. Mozes, M., Kleinberg, B.: No intruder, no validity: Evaluation criteria for privacy-preserving text anonymization (2021) 2
13. Narayanan, A., Shmatikov, V.: Robust de-anonymization of large sparse datasets. In: 2008 IEEE Symposium on Security and Privacy (sp 2008). pp. 111–125 (2008). <https://doi.org/10.1109/SP.2008.33> 5

14. Naveed, M., Kamara, S., Wright, C.V.: Inference attacks on property-preserving encrypted databases. In: Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security. pp. 644–655 (2015). <https://doi.org/10.1145/2810103.2813651> 4
15. Niedermeyer, F., Steinmetzer, S., Kroll, M., Schnell, R.: Cryptanalysis of basic bloom filters used for privacy preserving record linkage. *Journal of Privacy and Confidentiality* **6**(2) (2014). <https://doi.org/10.29012/jpc.v6i2.640> 4
16. Noever, D.: The Enron corpus: Where the email bodies are buried? (2020) 2
17. Osborne, J.D., O’Leary, T., Nadimpalli, A., Aly, S.M., Kennedy, R.E.: BRAT-synthetic: Text de-identification using a markov chain replacement strategy for surrogate personal identifying information (2022) 2
18. Packhäuser, K., Gündel, S., Münster, N., Syben, C., Christlein, V., Maier, A.: Deep learning-based patient re-identification is able to exploit the biometric nature of medical chest X-ray data. *Scientific Reports* **12**(1), 14851 (2022). <https://doi.org/10.1038/s41598-022-19045-3> 1
19. Packhäuser, K., Gündel, S., Thamm, F., Denzinger, F., Maier, A.: Deep learning-based anonymization of chest radiographs: A utility-preserving measure for patient privacy. In: *Lecture Notes in Computer Science*. pp. 262–272 (2023). https://doi.org/10.1007/978-3-031-43898-1_26 1
20. Peñas, A., Rodrigo, A.: A simple measure to assess non-response. In: Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics. pp. 1415–1424 (2011), <https://aclanthology.org/P11-1142/> 5
21. Pilán, I., Lison, P., Øvrelid, L., Papadopoulou, A., Sánchez, D., Batet, M.: The text anonymization benchmark (TAB): A dedicated corpus and evaluation framework for text anonymization. *Computational Linguistics* **48**(4), 1053–1101 (2022). https://doi.org/10.1162/coli_a_00458 2
22. Rocher, L., Hendrickx, J.M., de Montjoye, Y.A.: Estimating the success of re-identifications in incomplete datasets using generative models. *Nature Communications* **10**(1) (2019). <https://doi.org/10.1038/s41467-019-10933-3> 5
23. Tayebi Arasteh, S., Arias-Vergara, T., Pérez-Toro, P.A., Weise, T., Packhäuser, K., Schuster, M., Noeth, E., Maier, A., Yang, S.H.: Addressing challenges in speaker anonymization to maintain utility while ensuring privacy of pathological speech. *Communications Medicine* **4** (2024). <https://doi.org/10.1038/s43856-024-00609-5> 1
24. Trienes, J., Joseph, S., Schlötterer, J., Seifert, C., Lo, K., Xu, W., Wallace, B.C., Li, J.J.: InfoLossQA: Characterizing and recovering information loss in text simplification. In: Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). pp. 4263–4294. Association for Computational Linguistics, Bangkok, Thailand (2024). <https://doi.org/10.18653/v1/2024.acl-long.234> 2
25. Uppala, R.: OpenAI privacy filter: A cross-lingual, cross-domain PII evaluation across 32 benchmarks (2026) 2
26. Vakili, T., Dalianis, H.: Utility preservation of clinical text after de-identification. In: Proceedings of the 21st Workshop on Biomedical Language Processing. pp. 383–388. Association for Computational Linguistics (2022). <https://doi.org/10.18653/v1/2022.bionlp-1.38> 2
27. Vats, G., Agrawal, A., Singhal, S., Dash, A., Selvaraj, P., Jhavar, V., Chenna, R.P., Bharadwaj, Y.M.G.: REDACT: A systematically controlled multilingual benchmark for personal information detection (2026) 2

28. Vatsalan, D., Christen, P., O’Keefe, C.M., Verykios, V.S.: An evaluation framework for privacy-preserving record linkage. *Journal of Privacy and Confidentiality* **6**(1) (2014). <https://doi.org/10.29012/jpc.v6i1.636> 4
29. Vidanage, A., Ranbaduge, T., Christen, P., Schnell, R.: A taxonomy of attacks on privacy-preserving record linkage. *Journal of Privacy and Confidentiality* **12**(1) (2022). <https://doi.org/10.29012/jpc.764> 4